Florian Vahl, Dominik Buchhardt

1   The k-means algorithm

In the Initialization Phase, k random Points are generated, which represent the Centroids of the Clusters. In *digit_KMEANS.ipynb* the k is set to 10, the number of digits in base 10.

In the following Learning Phase the Data Points are assigned to the closest Centroid, the Centroids are reassigned to the means of their cluster. This Learning Phase repeats itself as long the Centroids still move (or in practice do not move more than a certain threshold as floating point numbers have limited accuracy on machines). The KMeans does not only find clusters, but the locally optimal cluster by minimizing the distance of each data point in a cluster to its centroid.

This Cluster Model can be used to assign new data points to the clusters. The *digit_KMEANS.ipynb* example assigns a digit image to the corresponding cluster which each represent a digit in base 10, thus 'recognizing' an image as a certain digit.

2   What is the best k?

3   Clustering comparisons

The K-Means algorithm creates two clusters on the nested circles with each cluster having half of each circle and their centroids being diametrically opposed on the inner circle. This result is not surprising as the K-Means algorithm finds clusters as the closest data points to a centroid, making the clusters circular and suitable for complex and possible nested clusters.

This behaviour can be also observed on the dataset with interlocked half moons, where each cluster claim part of the other half moon, creating incorrect categorizations. The Centroids are not on the half moons itself but on the hollow space of the moons.

In contrast to K-Means, the Agglomerative Clustering with single linkage (closest distance between two clusters) classifies the half moons correctly. This is not surprising as there is some empty space between each half moon, making the merging of clusters until two are left easy.

The Agglomerative Clustering with average (minimal average distance between two clusters) or complete (minimal maximum distance between two clusters) linkage yield the same result, with one cluster claiming a half moon and part of the other half moon. In this case it is worse than using single linkage because the number of false positives is higher, but it is still lower than using K-Means.

The DBScan algorithm yields for the half moon data set a perfect result, categorizing each half moon as its own cluster and when using more noise, classifying possible noise as noise.

In conclusion K-Means is more suited for classifying clusters with circular shapes, as circles have the minimal distance from their edges to their centers, making it unsuitable for datasets with non-circular/complex and possible nested clusters. Agglomerative Clustering merges clusters in a bottom-up approach in respect to certain distance criteria (like single, average or complete distance) between two clusters until a certain number of clusters are left. Using single linkage yields a result similar to using the DBScan algorithm for this dataset. DbScan uses a density-based approach of clustering, categorizing data points as core points, density-reachable from a core point and outliers, performing good with datasets with dense shapes/clusters even if they are complex. While DBscans performs better on these datasets, it's usage is more complex in *sklearn* than K-Means or Agglomerative Clustering, as the density needs to be set carefully to produce a satisfying result.

4   Self Organizing Maps